

Car Price Prediction with Linear Regression

This presentation explores a Python-based linear regression model. It predicts car prices using the scikit-learn library. Key steps include data preprocessing, model training, and performance evaluation.



by Ibrahim Abdel Nasser



Data Loading and Preprocessing

Loading the Dataset

First, import **pandas** to read the **car_prediction_data.csv** file.

```
alldata = pd.read_csv('car_prediction_data.csv')
```

<https://www.kaggle.com/datasets/zafarali27/car-price-prediction>

Rename and Drop Data

Drop unneeded data (Car ID).
Rename some of data (from Mileage to Kilometers).

```
dropped_alldata = alldata.drop(columns=["Model", "Car ID"])
```

Handling Missing Data

Then, fill missing values with the median. This ensures data completeness.

```
if data.isna().sum().sum() > 0:  
    print("Missing values detected, filling with median values.")  
    data.fillna(data.median(), inplace=True)  
else :  
    print(data.isna().sum())
```

Feature Engineering

Convert "Mileage" from miles to kilometers. Rename columns to improve readability.

```
data = dropped_alldata.rename(columns={"Year" : "Model Year", "Mileage" : "Kilometers"})  
data["Kilometers"] = data["Kilometers"] * 1.609
```

Puclic flew vorststure vcas553 Mearpain

Price of g:553 | ● Royal car



Exploratory Data Analysis

1 Kilometers vs. Price

A scatter plot visualizes the relationship. Price typically decreases as mileage increases.

```
plt.figure(figsize=(10, 6))
plt.scatter(data["Kilometers"], data["Price"], color='blue',
            marker='o', alpha=0.5)
plt.xlabel("Kilometers (KM)")
plt.ylabel("Price (EGP)")
plt.title("Mileage Vs Price")
plt.show()
```

2 Average Price per Brand

A bar chart displays median car prices. Each bar represents the median price per brand.

```
plt.bar(data["Brand"], data["Price"].median(),
        color='green')
plt.title("Average Car Price per Brand")
plt.xlabel("Model")
plt.ylabel("Price")
plt.show()
```


Data Splitting

1

Train-Test Split

`train_test_split` divides data.

2

80/20 Ratio

80% of data is for training and 20% is for testing.

3

Random State

`random_state=42` ensures reproducibility.

```
X = data.drop(columns=["Brand"])
y = data["Brand"]
X_train, X_test, y_train, y_test = train_test_split(X,
y, test_size=0.2, random_state=42)
```

TPAI Tetal

- 15 Itell05 Datal
- 12219,799-9880
- 754,00701922.2710

TTI THIITE

- 127len09: Saof
- 182:8.2909.5580
- 784,81102938,277)

TRINITG

TESTITG

TESITE

DATS/TICHETS

- 1.1etite fia1.39 Lead
- recertotion9,5od
- 1291,425 488,97.9690

CATEMCHETS

- 15tel05 Legnd
- futnry coters
- tall sapcy aty

Feature Selection and Encoding



Categorical columns include Brand, Fuel Type, Transmission, Condition, and Engine



LabelEncoder from scikit-learn converts categories to numerical labels.



Features are selected for model training after encoding is complete.

```
Size. categorical_cols = ["Brand", "Fuel  
Type", "Transmission", "Condition", "Engine Size"]  
label_encoders = {}  
for col in categorical_cols:  
    le = LabelEncoder()  
    data[col] = le.fit_transform(data[col])  
    label_encoders[col] = le
```


Model Training

1

Linear Regression

A **LinearRegression** model is initialized from **sklearn.linear_model**.

2

Model Fitting

The model learns from training data. It finds coefficients.

3

Prediction


The model predicts on the test set.

```
model = LinearRegression()  
model.fit(X_train, y_train)  
y_pred = model.predict(X_test)
```

Model Evaluation

Mean Squared Error

MSE measures prediction accuracy.



R-squared

R² indicates model fit. Ranges from 0 to 1.

```
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)
print(f"Mean Squared Error (MSE): {mse:.2f}")
print(f"R-squared (R2): {r2:.4f}")
```

Summary and Next Steps

Key Takeaways

Linear regression predicts car prices. **Scikit-learn** simplifies the process.

Further Exploration

Improve the model with more features. Consider other regression models.

